Problem 2: Neural Networks

Claudio Lopez, Diogo Silva

Universidad de Chile

October 20, 2024

Problem 2 overview

Goals

- Build a neural network capable of diagnosing breast cancer (binary classification)
- Define a loss function
- Find the optimal network parameters: activation function, optimizer, number of layers and number of neurons
- Investigate the effect of normalizing the training data
- Build a classifier for a multiple class problem

Exercise 1: Introducing Log-Likelihood

Maximum Likelihood Estimation

$$\max_{\theta} \prod_{i=1}^{N} p_\theta(y^{(i)}|x^{(i)})$$

where for each θ,

$$p_\theta(y|x) \geq 0,$$

$$\sum_{i=1}^{N} p_\theta(y^{(i)}|x^{(i)}) = 1$$

Exercise 1: Introducing Log-Likelihood

Finding optimal θ

$$\hat{\theta} = \arg\max_{\theta} \prod_{i=1}^{N} p_\theta(y^{(i)}|x^{(i)})$$

$$= \arg\max_{\theta} \frac{1}{N} \sum_{i=1}^{N} \log\left(p_\theta(y^{(i)}|x^{(i)})\right)$$

Exercise 2: Defining the Loss Function

Cross-entropy

Given two discrete probability p and q with the same support χ, the **cross-entropy** of p and q is given by

$$H(p, q) = -\sum_{x \in \chi} p(x) \log(q(x))$$

Cross-entropy loss function

The **cross-entropy** loss function is defined as

$$L(q, p) = \frac{1}{N} \sum_{i=1}^{N} H(p_i, q_i)$$

Exercise 2: Defining the Loss Function

Finding optimal θ

$$\begin{aligned}\arg\min_{\theta} L(p_\theta, p) &= \arg\min_{\theta} \frac{1}{N}\sum_{i=1}^{N} H(p_i, p_{\theta_i}) \\ &= \arg\max_{\theta} \frac{1}{N}\sum_{i=1}^{N} \log(p_\theta(y^{(i)}|x^{(i)})) = \hat{\theta}\end{aligned}$$

Metrics for the Classifier

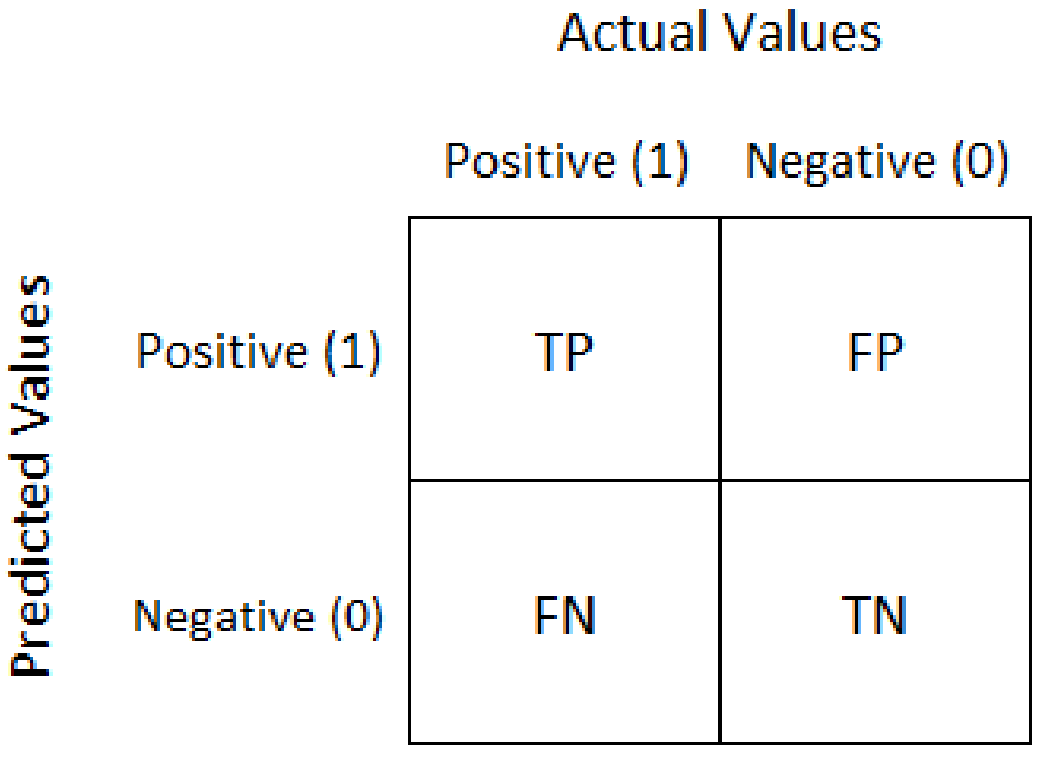

Metrics for the Classifier

Precision

Precision is the ratio of true positive predictions to the total positive predictions made by the model. It is defined as:

$$\text{Precision} = \frac{TP}{TP + FP}$$

Recall

Recall measures the ratio of true positive predictions to the total actual positives. It is defined as:

$$\text{Recall} = \frac{TP}{TP + FN}$$

Metrics for the Classifier

Matthews Correlation Coefficient (MCC)

Matthews Correlation Coefficient is a more informative metric that takes into account all four categories of the confusion matrix. It is defined as:

$$\text{MCC} = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}$$

MCC returns a value between -1 and 1, where 1 indicates a perfect prediction, 0 indicates no better than random prediction, and -1 indicates total disagreement between prediction and observation.

Exercise 3: Testing the first Neural Network

Network 1 Parameters

- Number of Layers: 2
- Number of Neurons: 16
- Activation Function: Sigmoid $:= \frac{1}{1+e^{-x}}$
- Optimizer: SGD

Exercise 3: Testing the First Neural Network

Average Network 1 Performance

- **Precision:** 0.6833
- **Recall:** 0.9938
- **MCC:** 0.1733

Average Network 1 Confusion Matrix

	Predicted Positive	**Predicted Negative**
Actual Positive	71.56	0.44
Actual Negative	35.16	6.84

Exercise 4: Testing Different Activation Functions

Network 2 Parameters

- Number of Layers: 2
- Number of Neurons: 16
- Activation Function: ReLU $:= \max(0, x)$
- Optimizer: SGD

Network 3 Parameters

- Number of Layers: 2
- Number of Neurons: 16
- Activation Function: Tanh $:= \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- Optimizer: SGD

Exercise 4: Testing Different Activation Functions

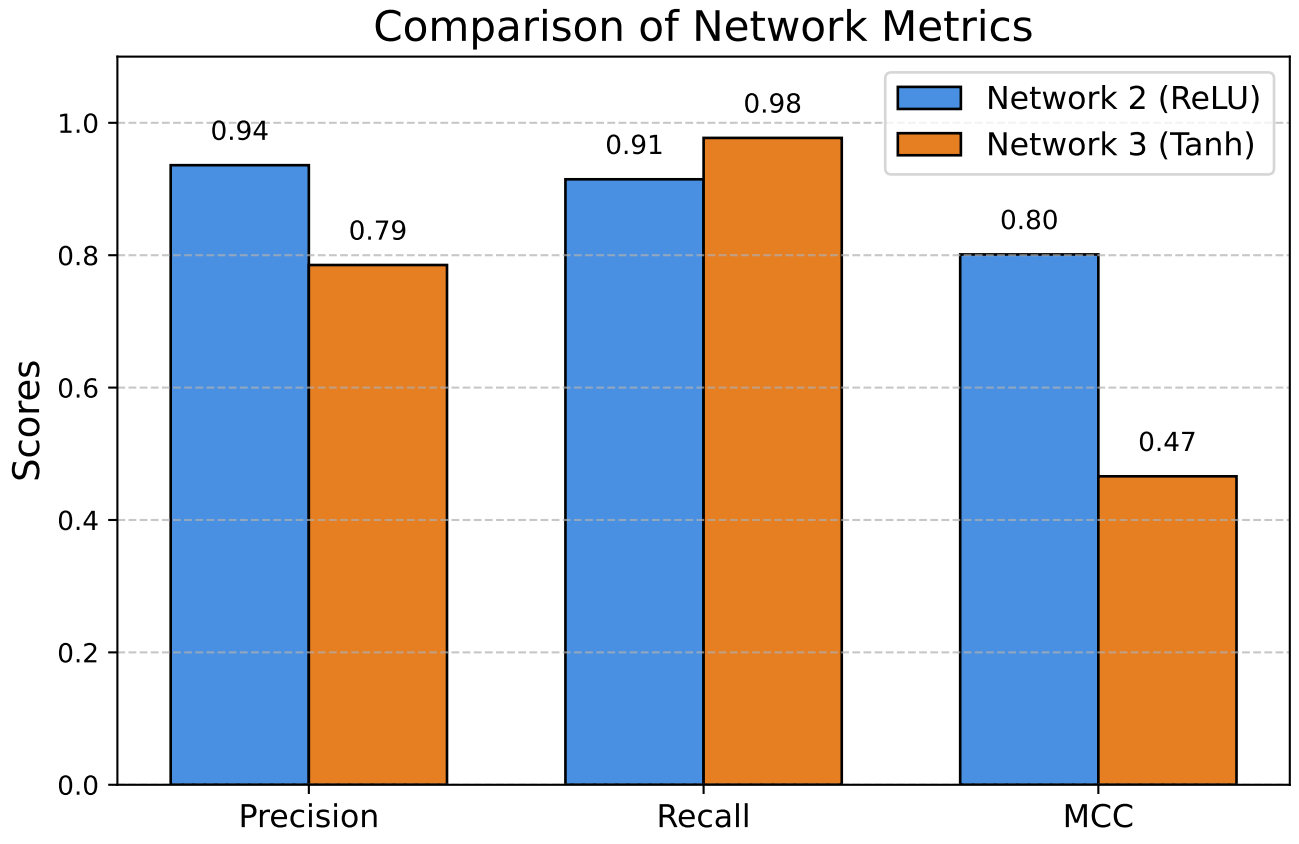

Exercise 4: Testing Different Optimizers

Network 4 Parameters

- Number of Layers: 2
- Number of Neurons: 16
- Activation Function: Sigmoid
- Optimizer: Adam

Network 5 Parameters

- Number of Layers: 2
- Number of Neurons: 16
- Activation Function: Sigmoid
- Optimizer: RMSprop

Exercise 4: Testing Different Optimizers

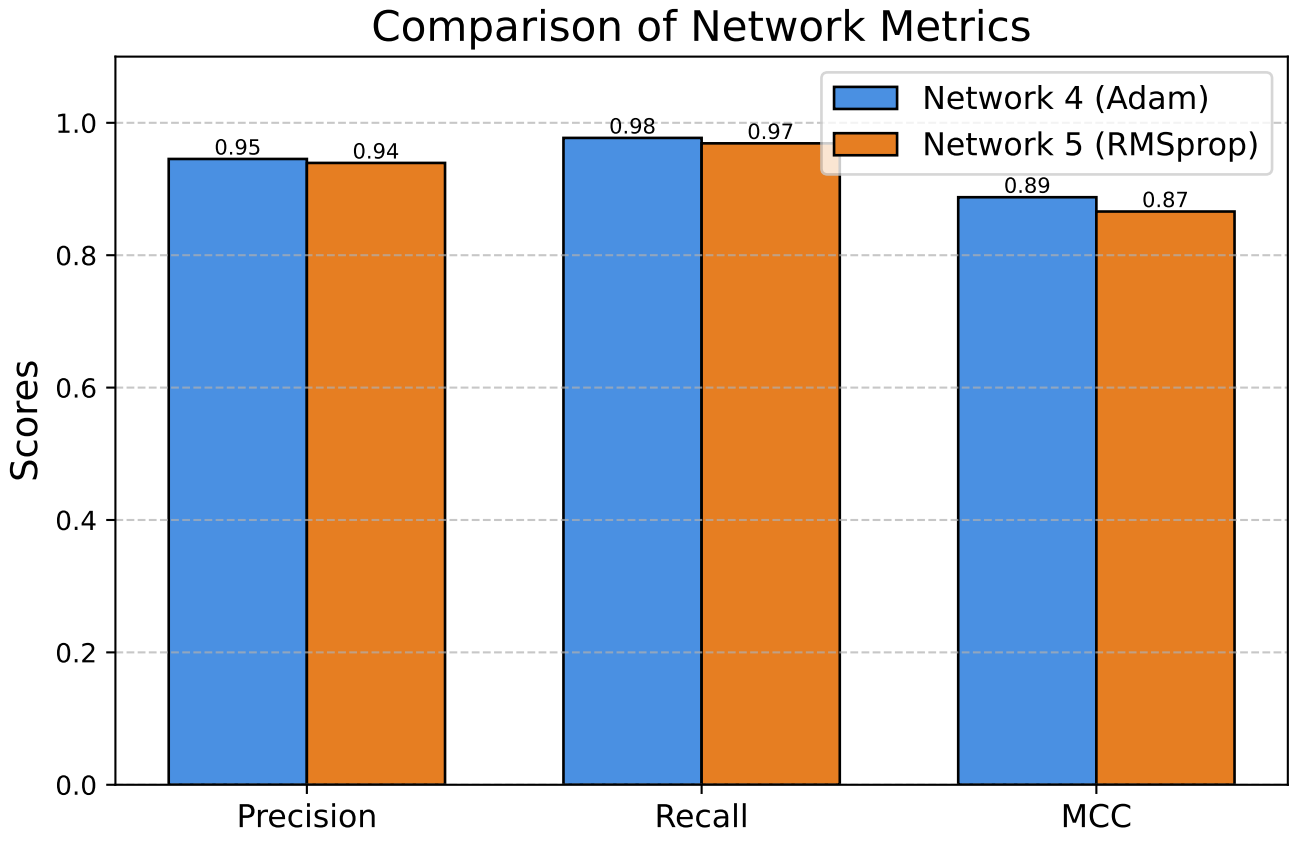

Exercise 5: Different Number of Layers and Neurons

Network 6 Parameters

- Number of Layers: 2
- Number of Neurons: 64
- Activation Function: ReLU
- Optimizer: Adam

Network 7 Parameters

- Number of Layers: 8
- Number of Neurons: 16

Network 8 Parameters

- Number of Layers: 4
- Number of Neurons: 32

Exercise 5: Different Number of Layers and Neurons

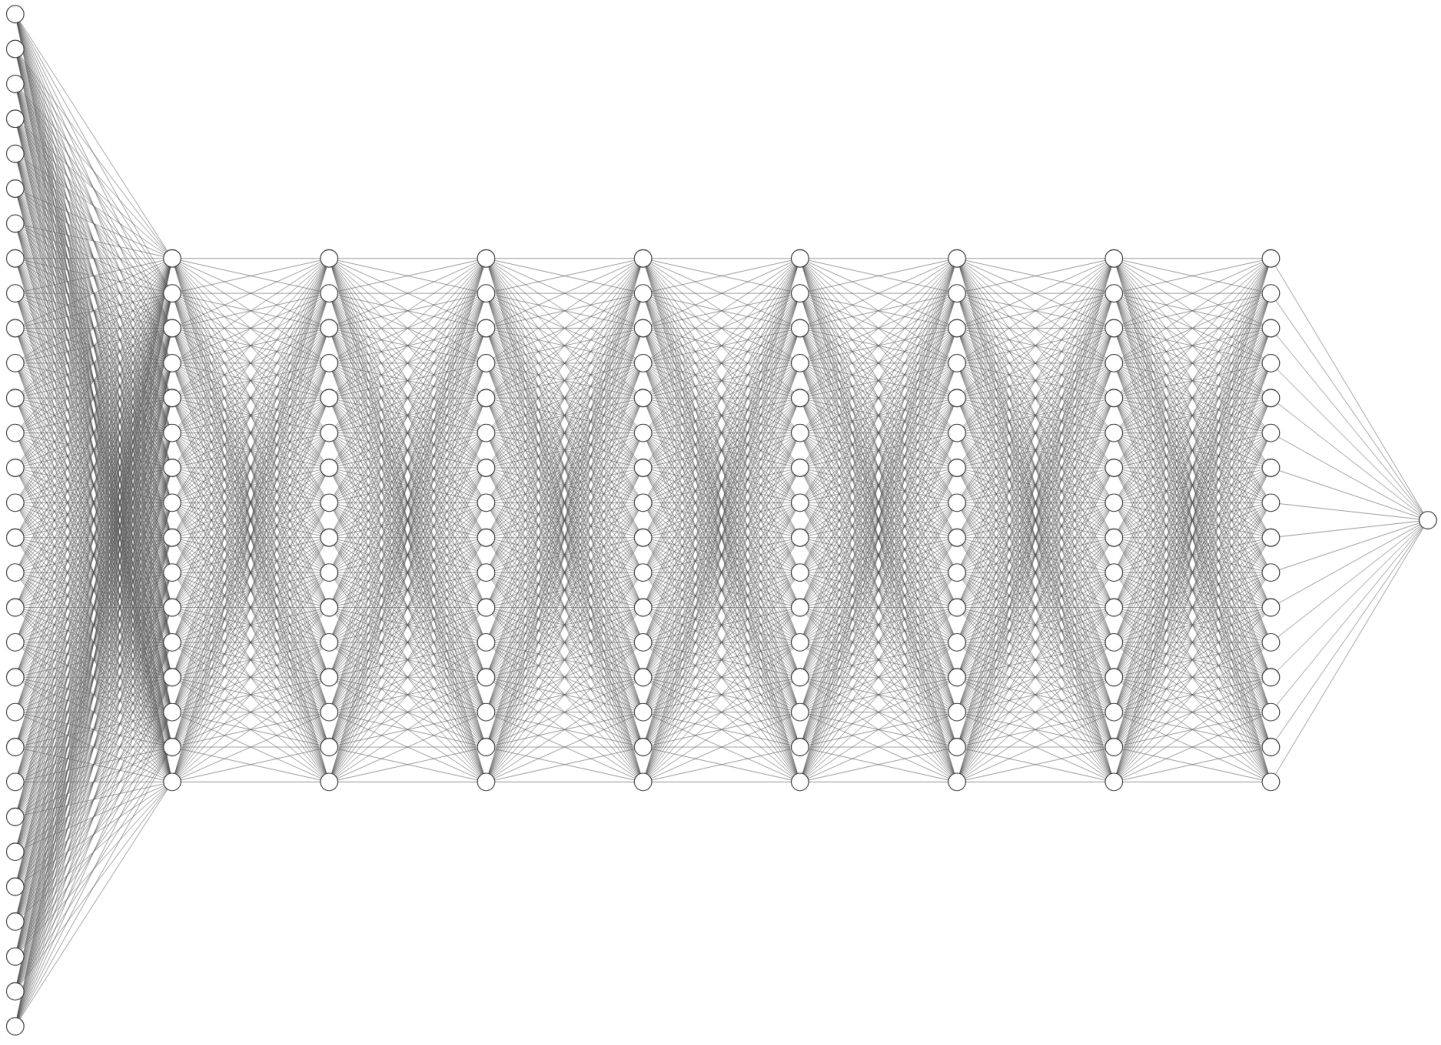

Figure: Network 7: 8 Layers, 16 Neurons

Exercise 5: Different Number of Layers and Neurons

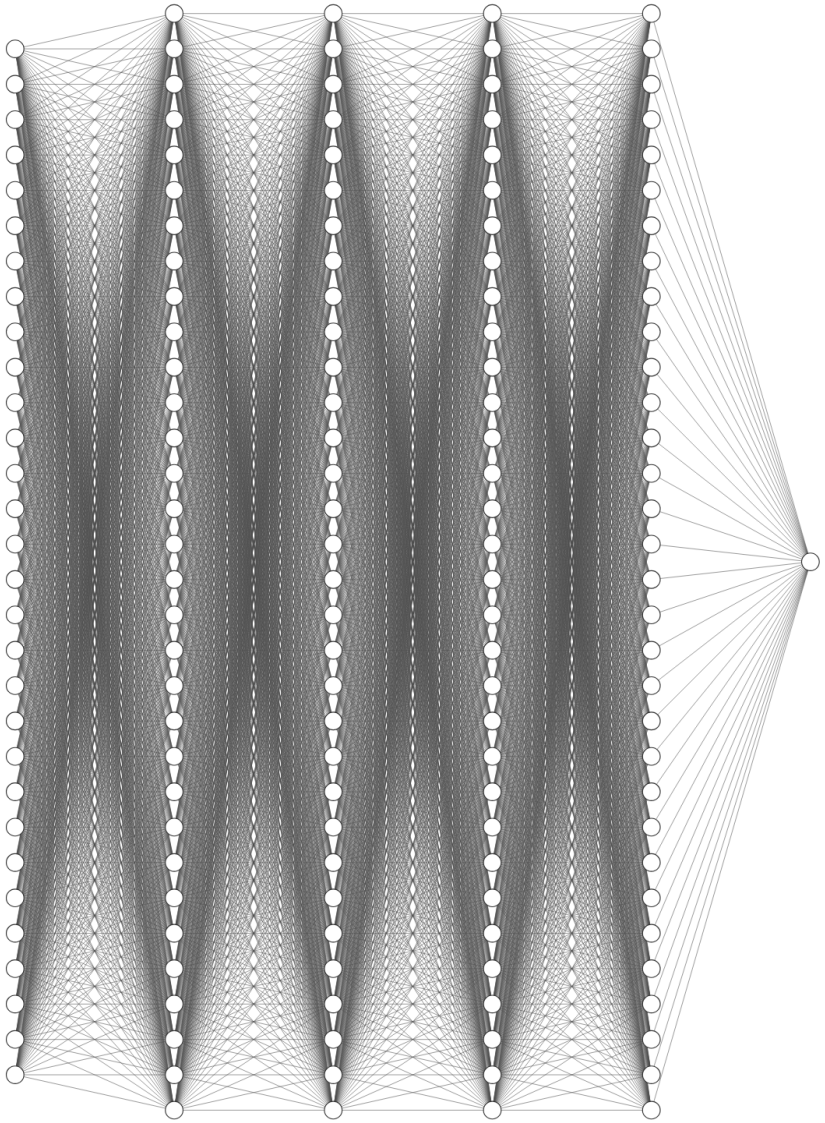

Figure: Network 8: 4 Layers, 32 Neurons

Exercise 5: Different Number of Layers and Neurons

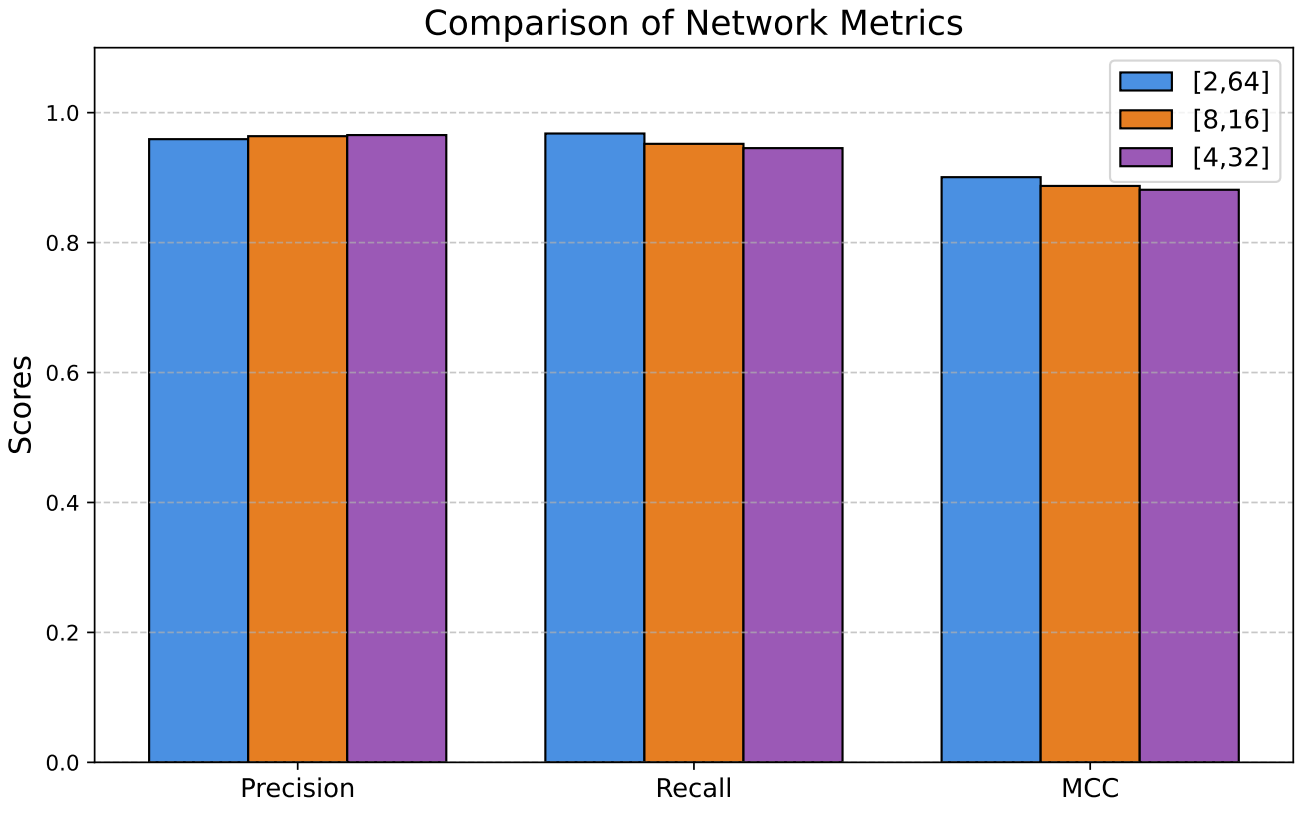

Exercise 6: Normalizing the Data

Network 9 Parameters

- Training and Testing sets normalized
- Number of Layers: 2
- Number of Neurons: 64
- Activation Function: ReLU
- Optimizer: Adam

Exercise 6: Normalizing the Data